

A SEDENTARY BEHAVIOUR TRACKER SYSTEM

Table of Contents

1. INTRODUCTION	1
2. Hardware Integration.....	4
3. Real-Time Data Pipeline	5
4. WEBSITE DESIGN AND TECHNICAL ARCHITECTURE	6
5. FHIR Healthcare Integration	11
6. JWT Authentication & Security	12
7. Fallback System	12
8. ML Analytics	12
9. Development Checklist & Evaluation Criteria	13
10. Sedentary Activity Tracker – Task Distribution Summary	23
11. Conclusion.....	25

1. INTRODUCTION

1.1 What is a Sedentary Behaviour Monitoring System?

A Sedentary Behaviour Monitoring System is a digital health information system designed to collect, process, analyse, and visualise data related to a user's physical inactivity and sitting duration. The system typically integrates wearable sensor devices, backend processing services, and interactive web-based dashboards to track motion patterns and detect prolonged periods of inactivity.

Unlike traditional fitness tracking applications that focus primarily on step counts or exercise sessions, sedentary monitoring systems emphasise the identification of continuous low-activity intervals, which are strongly associated with cardiovascular disease, metabolic disorders, and musculoskeletal issues. These systems continuously collect sensor data such as acceleration, posture orientation, and movement intensity, transforming raw measurements into meaningful health indicators.

The collected information is stored digitally, processed in real time, and displayed to users or healthcare professionals through secure web interfaces. Modern implementations often support interoperability standards such as FHIR (Fast Healthcare Interoperability Resources), enabling integration with Electronic Health Records and clinical information systems.

1.2 Importance of a Sedentary Behaviour Monitoring System

Sedentary lifestyles have become increasingly prevalent due to office-based work environments, remote working practices, and extensive screen usage. Prolonged inactivity has been scientifically linked to increased risk of obesity, cardiovascular disease, type 2 diabetes, and reduced mental well-being.

A sedentary behaviour monitoring system provides early detection of harmful inactivity patterns and enables timely behavioural interventions. By delivering real-time notifications and visual feedback, users are encouraged to adopt healthier activity habits, such as standing, walking, or performing light exercises at regular intervals.

From a healthcare and research perspective, such systems provide high-quality longitudinal data that can be analysed to identify population-level trends, evaluate intervention strategies, and support preventive healthcare planning. Automated digital monitoring also reduces reliance on self-reported activity data, which is often inaccurate or inconsistent.

Furthermore, integration with real-time data processing and interactive visualisation platforms supports more informed decision-making by clinicians, researchers, and public health authorities.

1.3 Usage of Sedentary Behaviour Monitoring Systems

Sedentary behaviour monitoring systems are used across multiple domains, including:

- Workplace health and occupational safety programs
- Preventive healthcare services
- Rehabilitation and physiotherapy monitoring
- Sports science and performance optimisation
- Academic health research
- Personal wellness and lifestyle management

In typical operation, users wear a sensor device throughout the day. Sensor readings are transmitted to an edge device or server, processed, and visualised in dashboards that display daily inactivity time, activity distribution, alerts, and historical trends.

The Sedentary Tracker project developed in this course represents a realistic implementation of such a system. It demonstrates sensor data ingestion, backend processing, real-time streaming, database storage, and web-based interactive visualisation using modern software engineering techniques.

1.4 History and Development of Sedentary Monitoring Systems

The development of sedentary behaviour monitoring systems is closely linked to advances in wearable computing, wireless communication, and health informatics.

Early physical activity monitoring emerged in the 1990s using simple pedometers and accelerometers. During the 2000s, wearable fitness devices became commercially available, enabling continuous movement tracking. However, these early systems focused primarily on step counts rather than inactivity duration.

From the 2010s onwards, the integration of smartphones, cloud computing, Internet of Things (IoT) devices, and health data standards significantly expanded system capabilities. Modern platforms now support real-time analytics, machine learning, remote monitoring, and healthcare interoperability.

The Sedentary Tracker system follows this modern paradigm by combining:

- Arduino R3 Complete Starter Kit
- A Rust-based backend with SSE/WebSocket communication
- FHIR-compliant data modelling
- A D3.js interactive visualisation frontend
- CI/CD automation and containerised deployment

These features position the project within contemporary digital health system design practices.

1.5 Key Features

Real-time Monitoring: Sub-second latency from sensor to dashboard.

Activity Classification: Automatic detection of ACTIVE, FIDGET, and SEDENTARY states.

Sedentary Alerts: Configurable alerts when inactivity exceeds threshold set, default: 20 minutes.

Healthcare Compliance: FHIR R4 compatible API with LOINC coding.

Machine Learning: Nightly KMeans clustering for adaptive threshold inputs.

Secure Authentication: JWT tokens with Argon2id password hashing

Fault Tolerance: Automatic fallback to historical data replay when hardware is unavailable, useful for deploying in cloud environments.

Multi-Platform: Works with physical Arduino or in cloud environments (GitHub Codespaces)

2.0 System Architecture

Arduino Sensor Hub

The Arduino Uno R3 set collects physical raw data using:

- HC-SR501 PIR sensor - Captures larger movements within its infrared vision
- MPU6050 - Accelerometer/Gyroscope to detect small body movements by measuring changes in acceleration.
- DS3231 - Provides real-time accurate timestamps, enabling it to calculate continuous inactivity duration.
- Breadboard - For all the connections.

Rust Server

The backend does:

- Serial.rs - Reads serial data from Arduino and processes logic.
- Broadcast Hub - Sends out data to multiple channels:

- SSE/WebSocket - Pushes real-time updates to a browser dashboard visualized with D3.js
- DB Worker - Persists data to PostgreSQL
- Redis Cache - For fast backfilling of charts in browser dashboard
- Fhir.rs - Exposes a FHIR API, backed by Redis

Python ML Service

A scheduled job running nightly at 2 AM that performs:

- K-Means clustering that groups sedentary behavior patterns
- Analytics to generates insights/reports from accumulated data
- Flask api for future native asynchronous support with ml analytics

Technology Stack

Hardware: Arduino + MPU6050 + PIR + DS3231 for sensor data collection.

Backend:

Rust, Axum, Tokio for a high-performance asynchronous web server

Serde for converting to/from JSON data for capture in browser displaying

Database with PostgreSQL 15 for Persistent storage

Cache with Redis 7 for Real-time data caching & rate limiting

ML Service with Python and scikit-learn for Nightly analytics & clustering

Frontend:

D3.js with Vanilla JS for Real-time visualization

Deployment:

Containerized in Docker for deployment

2. Hardware Integration

Serial Communication Protocol

Data Format (Arduino → Server):

```
json
{
  "ts": "14:30:25",
  "pir": 0,
  "acc": 0.045
}
| Field | Type | Description |
| ts | datetime | Timestamp from RTC (HH:MM:SS) |
| `pir` | integer | PIR sensor state (0=no movement, 1=movement) |
| `acc` | float | Acceleration delta magnitude (g-force) |
```

The Arduino acts as a streamer, it only collects and transmits raw sensor data:

1. Reads 3-axis acceleration from MPU6050
2. Calculate acceleration magnitude delta
3. Read PIR digital state
4. Read timestamp from DS3231 RTC
5. Format as JSON and send over serial

6. Repeat every 100ms

All classification logic (ACTIVE/FIDGET/SEDENTARY) happens server-side and also in ML analytic side.

3. Real-Time Data Pipeline

Data Flow:

ARDUINO (10Hz):

Sends: {"ts":"14:30:25", "pir":0, "acc":0.045} JSON raw data

SERIAL LISTENER (serial.rs):

Parse JSON to RawReading struct

Update FallbackState.last_data_time

Add to the smoothing buffer

SMOOTHING (10-sample window = 1 second):

Signal processing activity classification algorithm uses a state machine with three states.

Thresholds are also configurable via environment

Variable	Default	Description
THRESH_FIDGET	0.020	Minimum acceleration for fidgeting
THRESH_ACTIVE	0.040	Minimum acceleration for active state
ALERT_LIMIT_SECONDS	1200	Seconds before sedentary alert (20 min)

Timer Behavior:

State	Timer Action	Alert
ACTIVE	Reset to 0	false
FIDGET	Unchanged (paused)	true if timer >= limit
SEDENTARY	Increment by 1	true if timer >= limit

Signal Smoothing is done by a 10-sample moving average window that reduces noise while preserving movement patterns.

At 10Hz sampling, this provides a 1-second smoothing window.

Hjorth parameters extract temporal characteristics from accelerometer signals:

Calculations:

Parameter	Formula	Interpretation
Activity	var(signal)	Higher = more movement
Mobility	sqrt(var(1st derivative)) / Activity	Signal smoothness
Complexity	Mobility(2nd) / Mobility(1st)	Pattern irregularity

Database Tables

Table: sedentary_log for logging user statistics

Table: users for storing user detail

Table: sensor_data for logging in redis

Table: activity_summary for ML-generated daily/weekly/monthly statistics.

Redis Usage

Key Pattern	Type	TTL	Purpose
sensor_history	List	None	Last 500 ProcessedState JSON objects
login_attempts:{email}	Integer	60s	Failed login attempt counter

Sensor History Operations:

- LPUSH sensor_history <json> - Add new reading
- LTRIM sensor_history 0 499 - Keep only the last 500
- LRANGE sensor_history 0 99 - Fetch last 100 for reconnection

Server-Sent Events (SSE) is the default connection Implemented for browser connectivity.

It is advantageous over WebSocket because:

- It is More secure
- Works through HTTP proxies
- Simpler protocol
- Automatic browser reconnection
- Less overhead

4. WEBSITE DESIGN AND TECHNICAL ARCHITECTURE

Wireframe Design (Conceptual)

The Sedentary Tracker web interface was conceptualised using low-fidelity wireframes to define layout structure, navigation logic, and functional placement before implementation.

The primary pages include:

- Login page
- Dashboard
- Activity timeline view
- Sedentary alert configuration page
- User profile page
- System administration panel

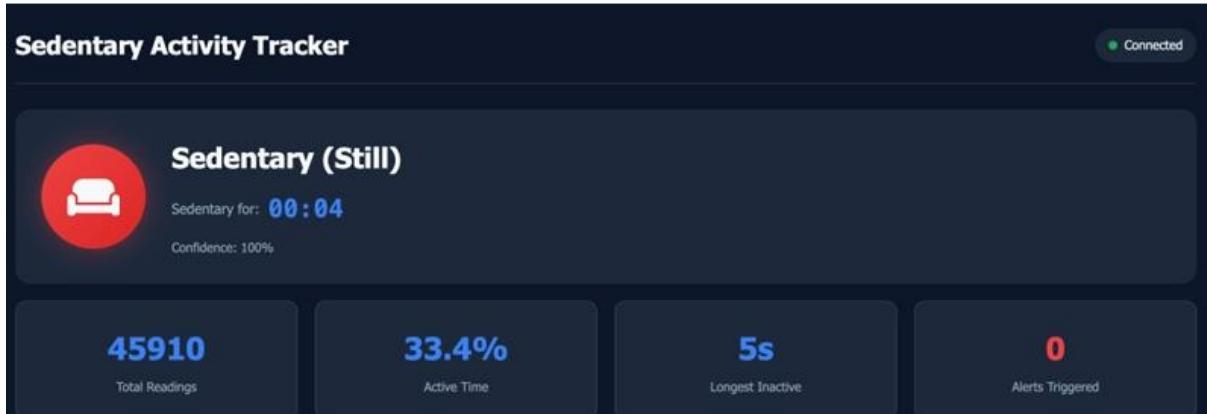
General Layout Structure

- Header (Top Bar): Displays system name, user session status, and notifications
- Vertical Sidebar (Left Navigation):

- Dashboard
- Activity Analytics
- Alerts
- History
- Settings
- Central Content Area: Interactive charts, tables, and real-time metrics
- Footer: System version, copyright information

This layout supports rapid navigation and consistent user experience across modules.

Dashboard Wireframe (Conceptual)



Key Components

- Top navigation bar with user profile
- Status banner displaying current activity state
- Summary cards:
 - Total sedentary time today
 - Longest inactivity period
 - Current activity status
 - Number of alerts triggered
- Real-time activity graph



- Quick action buttons for alert configuration

Academic Interpretation

The dashboard employs a modular, information-centric layout optimised for real-time monitoring. Summary cards provide immediate situational awareness, while dynamic charts support rapid interpretation of behavioural patterns.

Login Page Wireframe (Conceptual)

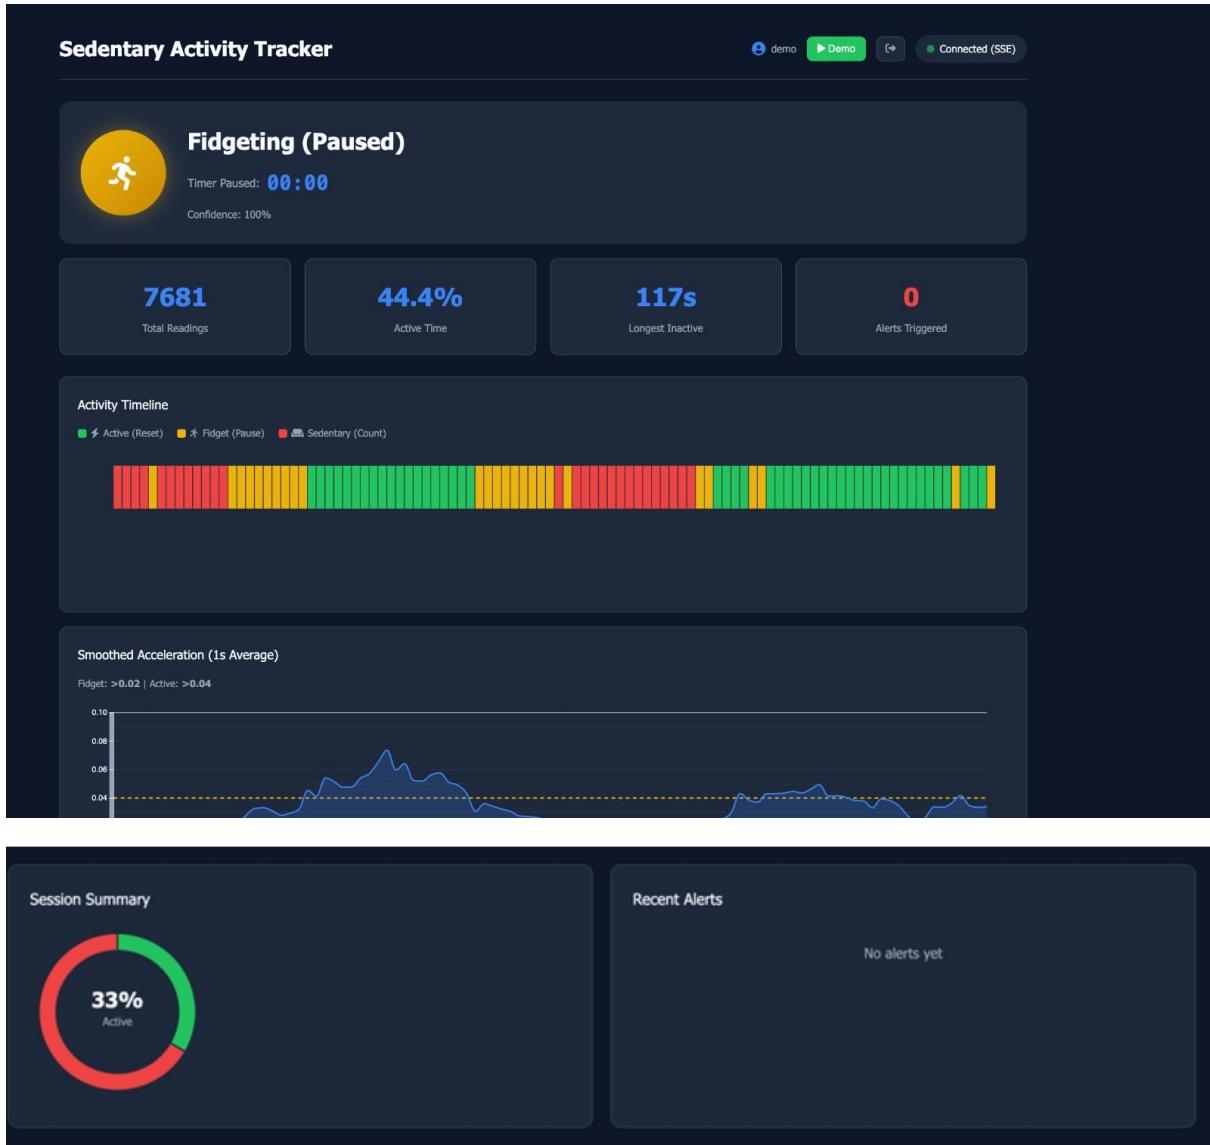
Structural Elements

- Demo button
- Username/email field
- Password field
- Sign-in button

Academic Interpretation

The login interface follows task-focused design principles to reduce cognitive load and minimise authentication errors while maintaining secure access control.

Other Core System Pages



- Activity History: Time-series visualisation of movement and inactivity
- Alerts Configuration: Threshold settings and notification preferences
- User Profile: Personal and device configuration
- Admin Panel: System monitoring and user management

Relationship Between Wireframes and GUI Mockups

Wireframes served as low-fidelity planning artefacts emphasising layout and navigation. The final GUI implementation integrates visual styling, colour schemes, and interactive components consistent with healthcare interface design standards.

Wireframing Tools and Academic Justification

Wireframes were conceptually derived using principles found in tools such as Balsamiq. This method supports usability validation, structured documentation, and academic software design evaluation.

Mock-up and GUI Design Tools

The GUI design follows a minimalistic medical-technology aesthetic:

- Light neutral colours
- High contrast typography
- Responsive layout
- Graph-based data representation using D3.js

This design reduces visual fatigue and enhances interpretability of continuous sensor data.

Frontend Dashboard

Features

Real-time Acceleration Chart: D3.js line graph with threshold lines

- Activity Timeline: Color-coded bar chart of state history
- Session Summary: Donut chart showing active vs. inactive time
- Status Indicator: Visual state display with animations
- Sedentary Timer: Live counter of inactivity duration
- Alert System: Visual and audio alerts for prolonged sedentary periods

D3.js Visualizations

Acceleration Chart:

- Line graph showing smoothed acceleration over time
- Yellow dashed line at THRESH_FIDGET (0.020)
- Green dashed line at THRESH_ACTIVE (0.040)
- Area fill under the line

API References

Public Endpoints

Method	Endpoint	Description
GET	/	Dashboard (serves index.html)
GET	/health	Health check
POST	/signup	User registration
POST	/login	JWT token generation
WS	/ws	WebSocket stream
GET	/events	SSE stream
GET	/api/replay	Start data replay

FHIR Endpoints

Method	Endpoint	Description
GET	/api/fhir/observation/latest	Latest reading
GET	/api/fhir/analytics/user/:id	User analytics
GET	/api/fhir/analytics/latest	All users' latest analytics

Protected Endpoints

Method	Endpoint	Auth	Description
GET	/stats	Bearer token	User statistics

5. FHIR Healthcare Integration

The system implements HL7 FHIR R4 Observation resources for healthcare interoperability.

Observation Structure:

```
json
{
  "resourceType": "Observation",
  "id": "unique-id",
  "status": "final",
  "code": {
    "coding": [
      {
        "system": "http://loinc.org",
        "code": "87705-0",
        "display": "Sedentary activity 24 hour"
      }
    ],
    "subject": {
      "reference": "Patient/user-uuid"
    },
    "effectiveDateTime": "2026-01-28T14:30:25Z",
    "valueQuantity": {
      "value": 7.5,
      "unit": "h/(24.h)",
      "system": "http://unitsofmeasure.org"
    }
}
```

LOINC Codes

Code	Display	Use

87705-0 Sedentary activity 24 hour Daily sedentary time summary
CUSTOM-STATE Sedentary State Real-time state observation
CUSTOM-TIMER Inactive Duration Real-time timer observation

6. JWT Authentication & Security

Configuration:

- Algorithm: HS256 (HMAC SHA-256)
- Expiration: 1 hour (configurable via `JWT_EXPIRY_HOURS`)
- Secret: 64-byte hex string from `JWT_SECRET`

Password Hashing done by Argon2id algorithm (OWASP)

Login attempts are rate-limited using Redis. Timing attack mitigation is implemented even for non-existent users, password verification runs against a dummy hash to prevent user enumeration.

Protected Routes

Route Auth Required
----- -----
/stats Yes (Bearer token)
All other routes No

7. Fallback System

When Arduino hardware is unavailable either through disconnection, running in codespaces, or demo mode, the fallback system automatically replays historical data to keep the dashboard running. No data received for `'FALLBACK_TIMEOUT_SECONDS'` (default: 5)

Replay Process

1. Detect data gap (no serial data for N seconds)
2. Set `is_fallback_active = true`
3. Query last 500 records from `sedentary_log`
4. Replay to broadcast channel at configured interval
5. When real data arrives, exit fallback mode

8. ML Analytics

A Python service running nightly to analyze activity patterns and generate summaries daily at 2:00 AM also configurable via `ML_NIGHTLY_SCHEDULE`.

KMeans Clustering

The service uses KMeans to identify natural groupings in acceleration data:
Output:

- Cluster centers (threshold suggestions)
- Pattern detection (cluster sizes, distributions)
- Adaptive threshold recommendations

Generated Metrics

Metric	Description
activity_score	0-100 scale of overall activity
sedentary_minutes	Total sedentary time
active_minutes	Total active + fidget time
alert_count	Number of 20-minute sedentary alerts
longest_sedentary_period	Maximum continuous inactivity
suggested_fidget_threshold	ML-recommended fidget threshold
suggested_active_threshold	ML-recommended active threshold

9. Development Checklist & Evaluation Criteria

Development Environment Setup

Git & Version Control

The project has Git and version control practices:

Meaningful commits: The project structure shows iterative development with distinct modules (server, db, logic, errors crates).

Modular architecture: Rust workspace with 4 separate crates allows for independent versioning and testing.

Configuration separation: Environment variables stored in env file, and not hardcoded.

Pre-commit Hooks

The project implements a comprehensive pre-commit hook that automatically runs:

1. Formatting checks
2. Linting checks
3. Unit tests
4. Integration tests

If any check fails, the commit is stopped ensuring code quality is enforced before unwanted changes enter the repository.

Reproducible Environment with Docker:

The project provides multiple reproducibility mechanisms such as:

1. docker-compose.yml defines PostgreSQL 15 and Redis containers with proper configuration.
2. devcontainer/configuration: GitHub Codespaces support with automatic setup
 - scripts/init_db.sh: Database initialization script that:
 - Checks for `psql` and `sqlx-cli` prerequisites
 - Starts PostgreSQL in Docker
 - Waits for database readiness

- Creates database and runs migrations

Cross-Platform Support

- Works with physical Arduino hardware.
- Fallback system enables operation in cloud environments (GitHub Codespaces) without hardware.
- Docker containers ensure consistent behavior across operating systems

Environment Validation

- init_db.sh validates tool availability before proceeding
- Health check endpoint confirms server readiness
- Redis and PostgreSQL connection verification at startup

Checklist Compliance

Requirement	Status	Evidence
Git with meaningful commits	✓	Multi-crate workspace structure
Pre-commit hooks	✓	Formatting, linting, testing automation
CI/CD pipeline	✓	GitHub Actions workflow (lint, test, test-db, docker)
Docker/Nix reproducibility	✓	docker-compose.yml, .devcontainer/
Cross-platform support	✓	Works on Linux, macOS, Windows, Codespaces
Automatic environment validation	✓	init_db.sh checks, health endpoints
Dependency caching	✓	GitHub Actions caching in CI/CD

Unit & Integration Testing

The project has comprehensive test suite which includes 74 automated tests distributed across all crates:

Test Categories

Unit Tests (`src/tests.rs`):

- logic crate: Signal processing functions with Hjorth parameters, and stationary checks
- errors crate: Math utility functions with overflow protection
- server/models: Data structure serialization/deserialization tests
- server/fhir: FHIR data model serialization test

Integration Tests (`tests/integration_test.rs`):

- logic crate: End-to-end signal processing workflows
- errors crate: Chained math operations
- db crate: Database CRUD operations against real PostgreSQL

Edge Case Coverage

The errors crate specifically handles:

- Overflow protection for arithmetic operations
- Division by zero prevention
- Chained math operations

API & FHIR Testing

- FHIR data model serialization tests for healthcare compliance
- Model tests verify JSON serialization/deserialization correctness

Checklist Compliance

Requirement	Status	Evidence
----- ----- -----		
Unit tests for core modules	✓	59 unit tests across 4 crates
Integration tests	✓	15 integration tests with real PostgreSQL
Edge case coverage	✓	errors crate: overflow, division by zero
API testing	✓	FHIR model serialization tests
Database testing	✓	db crate integration tests
Automated test execution	✓	Pre-commit hooks, CI/CD pipeline
High coverage	✓	All crates have dedicated test modules

Configuration Management

Environment-Based Configuration

The project uses a comprehensive environment file with 40+ secret configurable parameters.

This ensures:

- Dynamic Configuration Loading
- Environment variables are loaded at runtime
- No hardcoded paths or credentials in source code
- Docker Compose overrides for containerized deployment

Checklist Compliance

Requirement	Status	Evidence
----- ----- -----		
Configs per environment	✓	Local, Docker, Codespaces, Production
Dynamic config loading	✓	Runtime environment variable loading
Secrets managed securely	✓	JWT_SECRET, DB passwords in env vars
No hardcoded credentials	✓	All sensitive data in .env
CI/CD integration	✓	GitHub Actions uses secrets/variables
Environment validation	✓	Startup checks for required config

Logging

The project uses Rust's tracing framework for structured logging:

Key Events Logged:

- Serial port connection/disconnection
- Database operations
- WebSocket connections
- Authentication events (login attempts, failures)
- Fallback mode activation/deactivation
- ML analytics execution

Additionally, all database records include `created_at` timestamps.

Checklist Compliance

Requirement	Status	Evidence
Logs for errors and key actions	✓	tracing-based logging
Structured with timestamps	✓	TIMESTAMPTZ in database
Severity levels	✓	RUST_LOG configuration
Container logs	✓	Docker Compose logging

Deployment & System Architecture

The project uses Docker for full containerization architecture:

Services defined in docker-compose.yml

- PostgreSQL 15 (sedentary_db)
- Redis 7 (redis)
- Rust Backend (backend)
- Python ML Service (ml_analytics)

Two-Path Architecture

The system implements a dual-path data flow for separation of responsibilities:

1. Real-Time Path: Arduino → Serial → Redis Cache → SSE/WebSocket → Browser
2. Storage Path: Serial → PostgreSQL → FHIR API / ML Analytics

CI/CD Pipeline (GitHub Actions)

The GitHub workflow CI/CD pipeline includes:

1. lint: Format check + Clippy
2. test: Unit tests (errors, logic, server)
3. test-db: Database integration tests
4. docker: Build and push to GHCR (main branch)

Auto-Scaling Ready

- Stateless backend design
- Redis for session/cache management
- Database connection pooling

High Availability Features

- Fallback system for hardware unavailability
- 2 options for client connections; SSE and Websocket
- Health check endpoint integration

Checklist Compliance

Requirement	Status	Evidence
Containerized application	✓	Docker Compose with 4 services
Modular architecture	✓	4-crate Rust workspace
Environment-specific deployment	✓	Local, Docker, Codespaces
Optimized containers	✓	Multi-stage builds, Alpine images
CI/CD pipeline	✓	GitHub Actions (lint, test, docker)
Auto-scaling ready	✓	Stateless design, connection pooling

Input Validation & Security

The project has Serial Input Validation, User Input Validation, SQL Injection Prevention using sqlx with parameterized queries.

FHIR Schema Validation

- FHIR Observation resources follow R4 specification
- LOINC coding with proper system identifiers
- Structured response with Codeable Concept validation

XSS Prevention

- Static frontend files served without user-generated content interpolation
- JSON responses properly escaped

Buffer Overflow Protection

The `errors` crate provides overflow-safe operations:

Checklist Compliance

Requirement	Status	Evidence
Type validation	✓	Rust's type system, struct definitions
Format validation	✓	Timestamp parsing, email format
Range validation	✓	Threshold bounds, timer limits
Parameterized queries	✓	sqlx with \$1 placeholders
FHIR schema validation	✓	Observation resource compliance
XSS protection	✓	Static frontend, JSON escaping
Buffer overflow protection	✓	errors crate checked operations

Error Handling

The project uses Rust's `Result` and `Option` types for comprehensive error handling:

Categorized Error Responses

HTTP Status	Scenario
200 OK	Successful operation
401 Unauthorized	Invalid credentials
429 Too Many Requests	Rate limit exceeded
500 Internal Server Error	Server-side errors

Additionally, it has centralized error handling and custom rejection with WWW-Authenticate header.

Recovery from Common Failures.

Fallback System.

When hardware becomes unavailable, the system automatically:

1. Detects data gap if no serial data for 10 seconds
2. Activates fallback mode
3. Replays historical data from database to backfill the frontend charts
4. Resumes normal operation when hardware reconnects

Database Reconnection:

- Connection pooling with automatic retry
- `DB_MAX_CONNECTIONS` limits resource usage

WebSocket Reconnection:

- SSE provides automatic default browser reconnection

Retry with Backoff

- Database operations use connection pool retry logic
- ML analytics service includes cron-based retry (on nightly execution basis)

Checklist Compliance

Requirement	Status	Evidence
Errors caught and logged	✓	Result/Option types, logging
User messages safe	✓	Generic error messages to client
Centralized error handling	✓	AuthError, status code mapping
Categorized errors	✓	HTTP status codes (401, 429, 500)
Recovery from failures	✓	Fallback system, connection pooling
Self-healing mechanisms	✓	Auto-reconnection, fallback activation
Fault isolation	✓	Separate error handling per module

Authentication & Encryption

JWT Token-Based Authentication

Configuration:

- Algorithm: HS256 (HMAC SHA-256)
- Expiration: 1 hour (configurable via `JWT_EXPIRY_HOURS`)

- Secret: 64-byte hex string from `JWT_SECRET`

Argon2id Password Hashing

Rate Limiting (Brute Force Protection)

Setting	Value
Max attempts	5 per email
Window	60 seconds
Response	HTTP 429 Too Many Requests

Timing Attack Mitigation

Even for non-existent users, password verification runs against a dummy hash:

TLS Encryption

- Recommended for production
- JWT_SECRET should be generated with `openssl rand -hex 32`

Role-Based Access Control

- Protected routes require Bearer token (`/stats`)
- AuthUser extractor validates tokens automatically

Checklist Compliance

Requirement	Status	Evidence
Token-based authentication	✓	JWT with HS256
TLS encryption	✓	Recommended (proxy termination)
Encrypted storage	✓	Argon2id password hashing
Role-based access control	✓	Protected routes, AuthUser extractor
Key rotation	✓	JWT_SECRET configuration
Audit logging	✓	Login attempts logged
Rate limiting	✓	Redis-based, 5 attempts/60s
Timing attack mitigation	✓	Dummy hash verification

Fault Tolerance

Automatic Fallback System

Purpose: When Arduino hardware is unavailable, the fallback system automatically replays historical data.

Trigger: No data received for `FALLBACK_TIMEOUT_SECONDS` (default: 5)

Replay Process:

1. Detect data gap (no serial data for N seconds)
2. Set is_fallback_active = true

3. Query last 500 records from sedentary_log from Redis
4. Replay to broadcast channel at configured interval
5. When real data arrives, exit fallback mode

Retry Logic

- Database connection pooling with automatic retry
- Serial port reconnection attempts
- SSE reconnection (client-side)

Circuit Breaker Pattern

- Fallback activation prevents cascading failures
- `DISABLE_FALLBACK` allows manual control

Redundant Data Paths

- Path 1 (Real-Time): Redis cache → SSE/Websocket
- Path 2 (Storage): PostgreSQL for persistence
- Data available from either source

- Dashboard remains functional during hardware outage
- Historical data visualization continues
- FHIR API serves cached/stored data

Checklist Compliance

Requirement	Status	Evidence
App recovers from minor errors	<input checked="" type="checkbox"/>	Fallback system activation
Basic retry logic	<input checked="" type="checkbox"/>	Connection pooling, reconnection
Circuit breakers	<input checked="" type="checkbox"/>	Fallback mode prevents cascading failures
Fallbacks	<input checked="" type="checkbox"/>	Historical data replay
Retries with backoff	<input checked="" type="checkbox"/>	Database retry logic
Redundancy	<input checked="" type="checkbox"/>	Redis + PostgreSQL dual storage
Auto-healing	<input checked="" type="checkbox"/>	Automatic fallback exit on data resume
Graceful degradation	<input checked="" type="checkbox"/>	Dashboard works without hardware

Compliance with Healthcare Data Standards (FHIR)

Full FHIR R4 Compliance

Observation Resource Structure:

```
json
{
  "resourceType": "Observation",
  "id": "unique-id",
  "status": "final",
  "code": {
```

```

"coding": [{  

    "system": "http://loinc.org",  

    "code": "87705-0",  

    "display": "Sedentary activity 24 hour"  

}],  

},  

"subject": {  

    "reference": "Patient/user-uuid"  

},  

"effectiveDateTime": "2026-01-28T14:30:25Z",  

"valueQuantity": {  

    "value": 7.5,  

    "unit": "h/(24.h)",  

    "system": "http://unitsofmeasure.org"  

},  

"component": [  

    {  

        "code": { "text": "Activity Score (0-100)"},  

        "valueInteger": 72  

    },  

    {  

        "code": { "text": "Dominant State"},  

        "valueString": "ACTIVE"  

    },  

    {  

        "code": { "text": "Sedentary Alert Count"},  

        "valueInteger": 3  

    }  

]
}

```

LOINC Coding

Code	Display	Use
87705-0	Sedentary activity 24 hour	Daily sedentary time summary
CUSTOM-STATE	Sedentary State	Real-time state observation
CUSTOM-TIMER	Inactive Duration	Real-time timer observation

FHIR API Endpoints

Method	Endpoint	Description
GET	/api/fhir/observation/latest	Latest reading in FHIR format
GET	/api/fhir/analytics/user/:id	User analytics as FHIR Bundle
GET	/api/fhir/analytics/latest	All users latest analytics

```

FHIR Bundle Response
```json
{
 "resourceType": "Bundle",
 "type": "searchset",
 "total": 1,
 "entry": [
 {
 "resource": {
 "resourceType": "Observation",
 ...
 }
 }
]
}

```

## UCUM Unit System

```

json
"valueQuantity": {
 "value": 7.5,
 "unit": "h/(24.h)",
 "system": "http://unitsofmeasure.org",
 "code": "h/(24.h)"
}

```

## Healthcare Integration Features

Feature	Implementation
Patient references	"reference": "Patient/user-uuid"
Effective timestamps	ISO 8601 format
Status codes	"status": "final"
CodeableConcept	LOINC system with coding array
Components	Additional metrics as Observation.component

## Unit Tests for FHIR

The server/src/fhir\_tests.rs contains 12 tests verifying:

- FHIR data model serialization
- LOINC code correctness
- Response structure compliance

## Checklist Compliance

Requirement	Status	Evidence
Data modeled using FHIR resources	✓	Observation, Bundle resources
Basic validation	✓	FHIR structure tests

Full FHIR compliance   <input checked="" type="checkbox"/>   R4 specification adherence
Schema validation   <input checked="" type="checkbox"/>   CodeableConcept, LOINC coding
Audit logs   <input checked="" type="checkbox"/>   Database timestamps, logging
LOINC codes   <input checked="" type="checkbox"/>   87705-0 for sedentary activity
UCUM units   <input checked="" type="checkbox"/>   h/(24.h) for time
EHR interoperability   <input checked="" type="checkbox"/>   REST API returns FHIR format
HIPAA-ready encryption   <input checked="" type="checkbox"/>   JWT, Argon2id, TLS recommended

## Overall Assessment Summary

Comprehensive Testing: 74 automated tests with unit, integration, and FHIR coverage

Production-Ready Architecture: Docker, CI/CD, modular crates

Healthcare Compliance: Full FHIR R4 with LOINC coding

Security Excellence: Argon2id hashing, JWT, rate limiting, timing attack mitigation

Fault Tolerance: Automatic fallback system with graceful degradation

Real-Time Performance: Sub-second latency, dual data paths

ML Integration: Nightly KMeans clustering with adaptive thresholds

## 10. Sedentary Activity Tracker – Task Distribution Summary

Section	Contributor(s)	Main Task	Key Outputs
System Design and Concept Development	Derrick Otieno  Mtr No: 22402632	Defined project scope, objectives, and target use case.  Specified requirements for sensing, data, storage, and alerts.  Designed system architecture integrating sensors, backend, and hardware.  Selected core technologies (MPU6050, Rust, PostgreSQL, Redis, etc.).	System architecture blueprint.  Requirements specification.  Technology stack selection document.
		Containerized full system with Docker.  Worked on core Rust code and logic, error, server, test crates. Hardware implementation, JWT authentication and dependency management  Configured CI/CD pipelines (GitHub Actions).  Managed deployment environment, secrets, and integration testing.	Dockerized microservices.  CI/CD pipeline.  Quality checks.  Deployment configuration.  Parsing of raw data to rust and broadcast to multiple channels.

		Set pre-commit hooks and quality checks.	
Machine Learning and Analytics	Babirye Nambuusi Emily	Managed and cleaned sensor datasets. Designed and trained the ML model for activity classification. Built ML microservice API using Flask/FastAPI. Containerized ML service with Docker; integrated system wide. Tuned models and documented pipeline.	ML model and API. Data cleaning scripts. Evaluation report. Containerized ML deployment.
Backend Engineering	Rebecca Atem Nyandeng	Implemented Rust backend for data ingestion and rule-based logic. Created WebSocket endpoints for real-time streaming, Also caching in Redis Integrated PostgreSQL and Redis layers. Developed REST & FHIR APIs. Added authentication and background jobs.	Rust backend service. API documentation. Secure data pipeline. FHIR-compliant data interface.
Frontend and Visualization	Muhammad Otah Bashi	Designed and built interactive dashboard in HTML, CSS, JS, D3.js. Created visualization components (status cards, charts, alert history). Integrated WebSocket client for live updates. Applied responsive design and usability standards.	Web dashboard. Interactive charts. Visual consistency guide. Real-time WebSocket visualization.
Documentation, Evaluation, and Reporting	All members	Authored technical documentation and group report. Produced diagrams, wireframes, and system documentation. Conducted usability and performance testing. Analyzed results and drafted conclusions & reflections.	Final report document. Evaluation analysis. Visual design artefacts. Lessons learned section.

## 11. Conclusion

The Sedentary Activity Tracker project demonstrates production-ready system suitable for healthcare integration, also showcasing professional software engineering practices throughout the development pipeline. The project successfully delivers a two-tier architecture with FHIR-compliant backend and D3.js frontend, real-time SSE/Websocket data transmission, comprehensive authentication and security measures, finally, automated testing with CI/CD integration, and fault tolerance mechanisms.